Distributed AI Resource Fabric

Coordinate AI work across your available computers, oversee active workloads and review capacity and recovery decisions.

Product overview

An AI-resource coordination asset for reviewing capacity, workload placement and recoverability.

Potential acquirers

- AI infrastructure vendors
- Compute operators
- Enterprise AI platform teams

Documented functions

- Discover usable computing capacity
- Review workload placement
- Reserve compatible resources
- Observe execution and recovery
- Retain measured qualification evidence

Available output types

Capacity records / Workload decisions / Execution and recovery evidence

Defined work. Separate written terms.

Seller asking reference / AS IS

EUR 500,000 · indicative USD 560,700

Identified edition: 0.19.0

Developed software · Conditional handover

Current scope and conditions

Seller-owned first-party rights plus specifically agreed transferable software/delivery artifacts; website displays and third-party ownership rights excluded from standalone software asset value.

• Target-environment suitability, deployment acceptance and any required independent qualification are agreed separately.

Conditions before handover

• Dependency manifest exists without an identified lockfile; reproducibility should be reviewed.

• A detailed component inventory and licensing report are required for the agreed delivery scope.

• Third-party licensing and obligations documentation remains incomplete.

• A security assessment or an explicit statement of its absence is required before acceptance.

• An explicit delivery manifest and handover scope remain to be agreed.

• Agree and freeze the exact unmerged v0.24 delivery scope; main remains v0.19 and no release promotion is implied.

• Physical multi-GPU model execution, two-host model execution/recovery, live kernel containment and production acceptance remain unqualified.

The seller offers only the expressly agreed first-party rights it owns or controls, to the extent transferable. Third-party materials retain their own terms. Independent title and freedom-to-operate opinions are not established.

Prices & sources.

Price basis: 2026-11-10. EUR is authoritative; market value is not established.

Price / model reference	EUR	USD - indicative
Seller asking reference	500,000	560,700
Negotiation reference - low	390,000	437,346
Negotiation reference - high	630,000	706,482
Recreation cost estimate	1,400,000	1,569,960
Marketability cost scenario	2,310,000	2,590,434
Monetisation cost scenario	3,080,000	3,453,912
Risk-adjusted strategic reference	3,721,323	4,173,092
Weighted economic reference	2,730,463	3,061,941

Cost basis

Some dated work-package assumptions are disclosed, but they do not reconstruct the complete retained recreation anchor. A complete scope-specific budget is required.

Coverage describes stated seller estimates only. The retained earlier anchor lacks a complete work breakdown; added scope does not validate it.

Source references

WIPO — Intellectual Property Valuation Basics — The cost method

Knab — Software developer freelance rates in 2026

Upwork — Software developer hiring costs

External publishers inform method and rate context, not individual asset valuations. Effort and model factors are seller assumptions.

EUR 1 = USD 1.1214; ECB 2026-11-10.

https://thearchitect-max.github.io/MyProjects/projects/distributed-ai-resource-fabric/#pricing

Individual source review.

Reviewed 10 November 2026. Current main documentation, source inventory and open candidate status; no fresh execution of private product tests.

Retained source evidence

v0.19 source report: exact-source Colab Tesla T4 GPU admission and 298/298 repository tests; report flags no vLLM/llama.cpp model inference, multi-host failover, RDMA or pooled VRAM.

Remaining limits

Single T4 admission is not two-host distributed inference/VRAM pooling; pending multi-host tests and real worker failover remain unqualified.

Development position

validation: Kaggle T4x2 GPU sharding, Colab pairing and live cgroup2 host campaign; docs: mandatory distributed GPU model execution gates (v0.24 375/375 baseline confirmed); v0.24: 375/375 exact-source + Tesla T4 PASS; live kernel pending; v0.23: 367/367 exact-source + Tesla T4 PASS; live cgroup kernel pending; v0.22: 351/351 exact-source + Tesla T4 PASS; live kernel containment pending; v0.21: exact-source 335/335 + Tesla T4 PASS; live cgroup kernel gate pending; v0.20: trusted cgroup pre-exec launch — exact-source software/T4 PASS; live kernel gate pending

Acquisition evidence profile

Technical scope: V — An AI-resource coordination asset for reviewing capacity, workload placement and recoverability. The offered edition and material boundaries govern review.

Evidence and validation: E4 — E3 plus repository-retained independent/external validation evidence identified. Detailed source-specific records are reviewed under agreed confidentiality; independent intended-use acceptance remains separate.

Rights and component obligations: R2 — R1 plus dependency/third-party disclosure basis identified. Review the agreed first-party rights, third-party obligations and delivery perimeter.

Commercial basis: C2 — Seller ask and reproducible economic reference established. Market value is not established.

Delivery and handover: T3 — T2 plus validation, build/release and dependency-disclosure evidence. The listed conditions and written delivery agreement govern acceptance.

Buyer information and next step

Quote TA-IP-022 and the project title. Agree the exact first-party edition, rights, delivery manifest and recipient acceptance separately.

https://thearchitect-max.github.io/MyProjects/projects/distributed-ai-resource-fabric/